

*Operating  
Systems:  
Internals  
and Design  
Principles*

# Chapter 7 Memory Management

Eighth Edition  
William Stallings


|         |   |
|---------|---|
| Frame   | A fixed-length block of main memory.  |
| Page    | A fixed-length block of data that resides in secondary memory (such as disk). A page of data may temporarily be copied into a frame of main memory.   |
| Segment | A variable-length block of data that resides in secondary memory. An entire segment may temporarily be copied into an available region of main memory (segmentation) or the segment may be divided into pages which can be individually copied into main memory (combined segmentation and paging). |

**Table 7.1**

## **Memory Management Terms**

# Memory Management Requirements


- Memory management is intended to satisfy the following requirements:
  - Relocation
  - Protection
  - Sharing
  - Logical organization
  - Physical organization

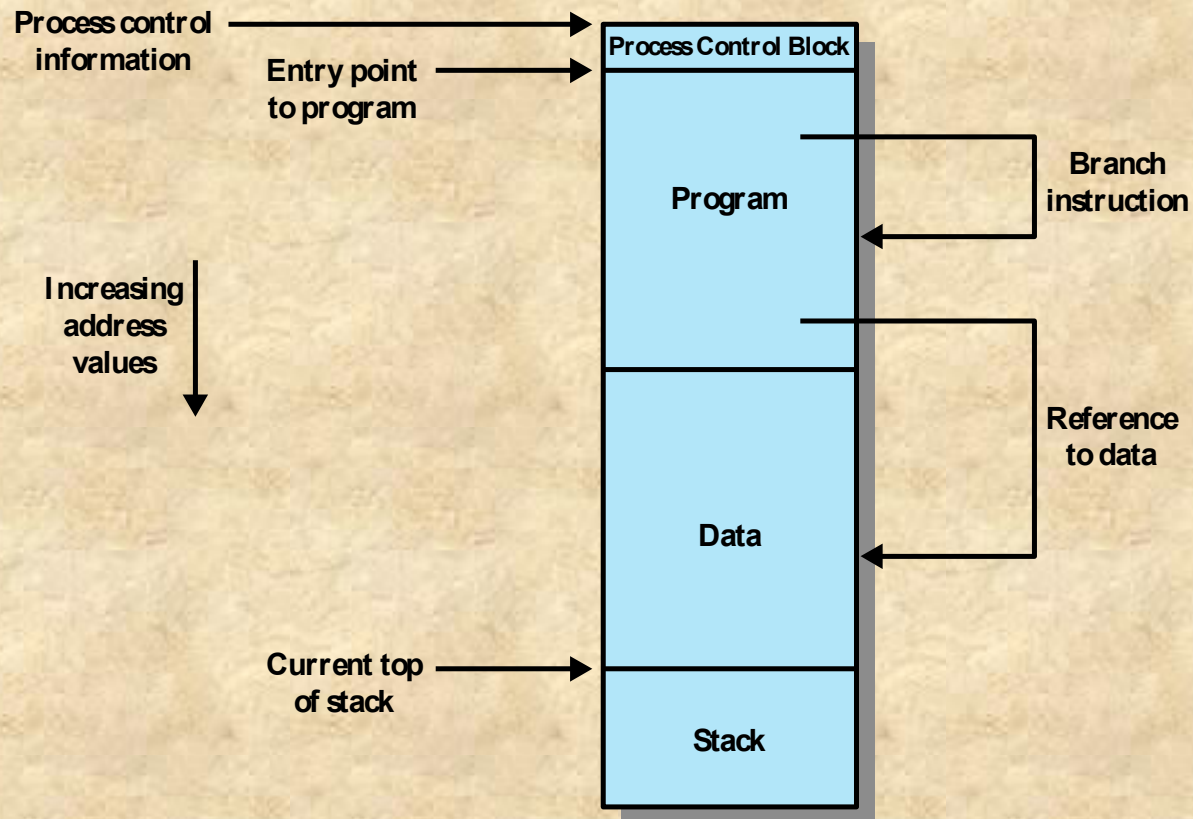




# Relocation

- Programmers typically do not know in advance which other programs will be resident in main memory at the time of execution of their program
- Active processes need to be able to be swapped in and out of main memory in order to maximize processor utilization
- Specifying that a process must be placed in the same memory region when it is swapped back in would be limiting
  - may need to *relocate* the process to a different area of memory





**Figure 7.1 Addressing Requirements for a Process**

# Protection


- Processes need to acquire permission to reference memory locations for reading or writing purposes
- Location of a program in main memory is unpredictable
- Memory references generated by a process must be checked at run time
- Mechanisms that support relocation also support protection





# Sharing

- Advantageous to allow each process access to the same copy of the program rather than have their own separate copy
- Memory management must allow controlled access to shared areas of memory without compromising protection
- Mechanisms used to support relocation support sharing capabilities



# Logical Organization

- Memory is organized as linear

## Programs are written in modules

- modules can be written and compiled independently
- different degrees of protection given to modules (read-only, execute-only)
- sharing on a module level corresponds to the user's way of viewing the problem

- Segmentation is the tool that most readily satisfies requirements



# Physical Organization

Cannot leave the programmer with the responsibility to manage memory


Memory available for a program plus its data may be insufficient

Programmer does not know how much space will be available

*overlaying* allows various modules to be assigned the same region of memory but is time consuming to program

# Memory Partitioning

- Memory management brings processes into main memory for execution by the processor
  - involves virtual memory
  - based on segmentation and paging
- Partitioning
  - used in several variations in some now-obsolete operating systems
  - does not involve virtual memory




| Technique                          | Description  | Strengths  | Weaknesses  |
|------------------------------------|--|--|---|
| <b>Fixed Partitioning</b>          | Main memory is divided into a number of static partitions at system generation time. A process may be loaded into a partition of equal or greater size.  | Simple to implement; little operating system overhead.   | Inefficient use of memory due to internal fragmentation; maximum number of active processes is fixed. |
| <b>Dynamic Partitioning</b>        | Partitions are created dynamically, so that each process is loaded into a partition of exactly the same size as that process.  | No internal fragmentation; more efficient use of main memory.  | Inefficient use of processor due to the need for compaction to counter external fragmentation.        |
| <b>Simple Paging</b>               | Main memory is divided into a number of equal-size frames. Each process is divided into a number of equal-size pages of the same length as frames. A process is loaded by loading all of its pages into available, not necessarily contiguous, frames. | No external fragmentation.   | A small amount of internal fragmentation.   |
| <b>Simple Segmentation</b>         | Each process is divided into a number of segments. A process is loaded by loading all of its segments into dynamic partitions that need not be contiguous.   | No internal fragmentation; improved memory utilization and reduced overhead compared to dynamic partitioning.              | External fragmentation.   |
| <b>Virtual Memory Paging</b>       | As with simple paging, except that it is not necessary to load all of the pages of a process. Nonresident pages that are needed are brought in later automatically.  | No external fragmentation; higher degree of multiprogramming; large virtual address space.                                 | Overhead of complex memory management.  |
| <b>Virtual Memory Segmentation</b> | As with simple segmentation, except that it is not necessary to load all of the segments of a process. Nonresident segments that are needed are brought in later automatically.  | No internal fragmentation, higher degree of multiprogramming; large virtual address space; protection and sharing support. | Overhead of complex memory management.  |

**Table 7.2**

**Memory Management Techniques**

(Table is on page 315 in textbook)





(a) Equal-size partitions




(b) Unequal-size partitions

Figure 7.2 Example of Fixed Partitioning of a 64-M byte Memory


# Disadvantages

- A program may be too big to fit in a partition
  - program needs to be designed with the use of overlays
- Main memory utilization is inefficient
  - any program, regardless of size, occupies an entire partition
  - *internal fragmentation*
    - wasted space due to the block of data loaded being smaller than the partition





(a) One process queue per partition



(b) Single queue

**Figure 7.3** Memory Assignment for Fixed Partitioning




# Disadvantages

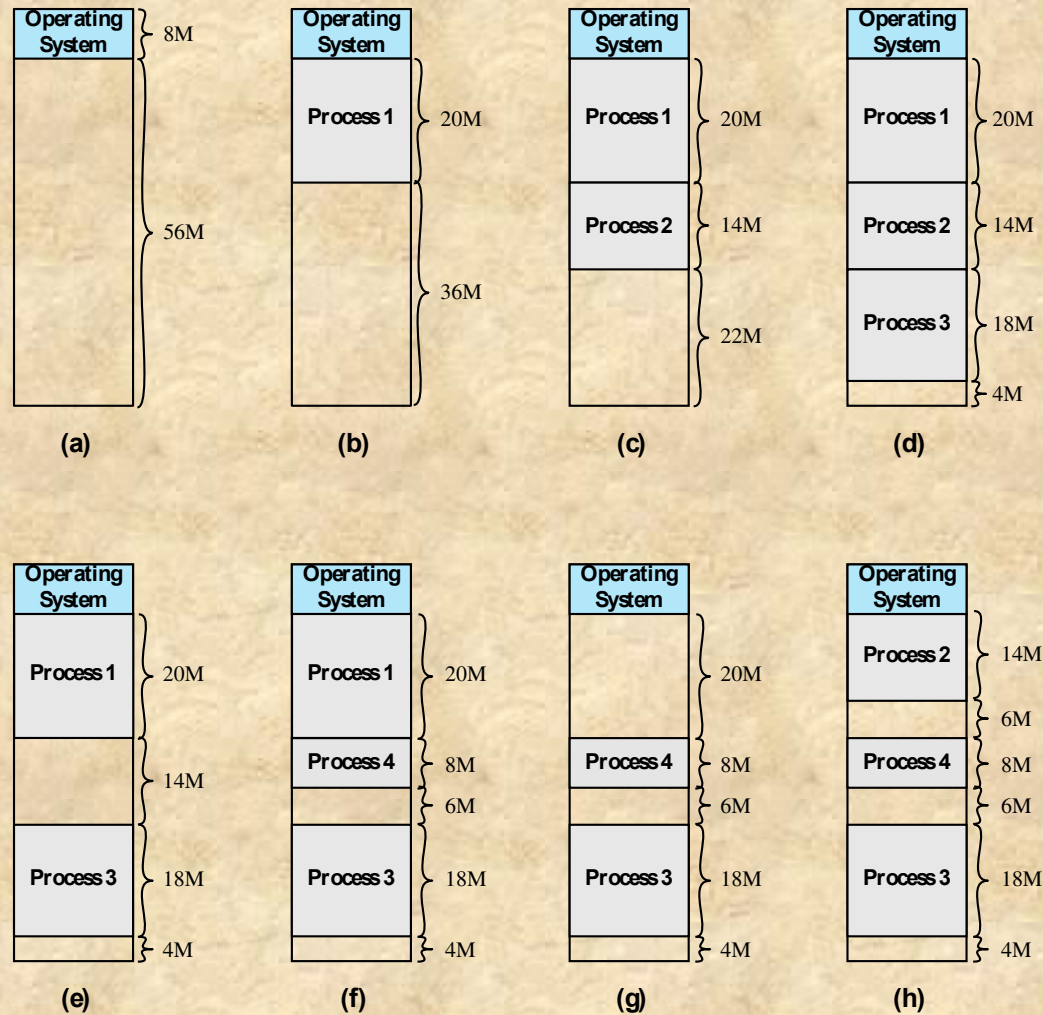
- The number of partitions specified at system generation time limits the number of active processes in the system
- Small jobs will not utilize partition space efficiently



# Dynamic Partitioning

- Partitions are of variable length and number
- Process is allocated exactly as much memory as it requires
- This technique was used by IBM's mainframe operating system, OS/MVT





**Figure 7.4 The Effect of Dynamic Partitioning**



# Dynamic Partitioning

## External Fragmentation

- memory becomes more and more fragmented
- memory utilization declines

## Compaction

- technique for overcoming external fragmentation
- OS shifts processes so that they are contiguous
- free memory is together in one block
- time consuming and wastes CPU time

# Placement Algorithms

## Best-fit

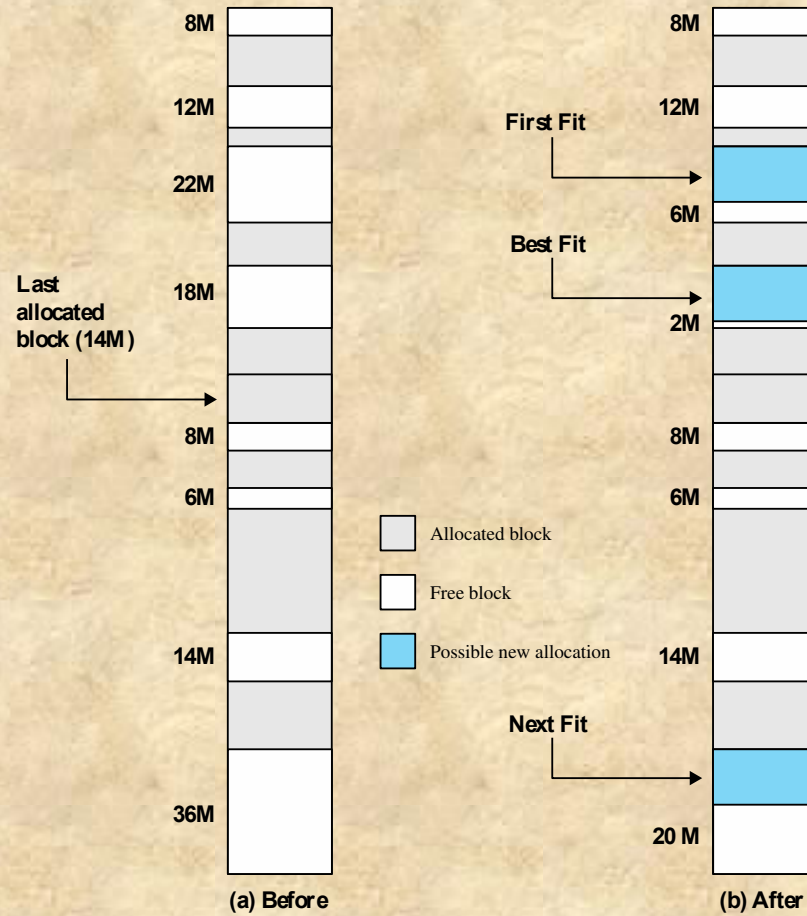
- chooses the block that is closest in size to the request

## First-fit

- begins to scan memory from the beginning and chooses the first available block that is large enough

## Next-fit

- begins to scan memory from the location of the last placement and chooses the next available block that is large enough



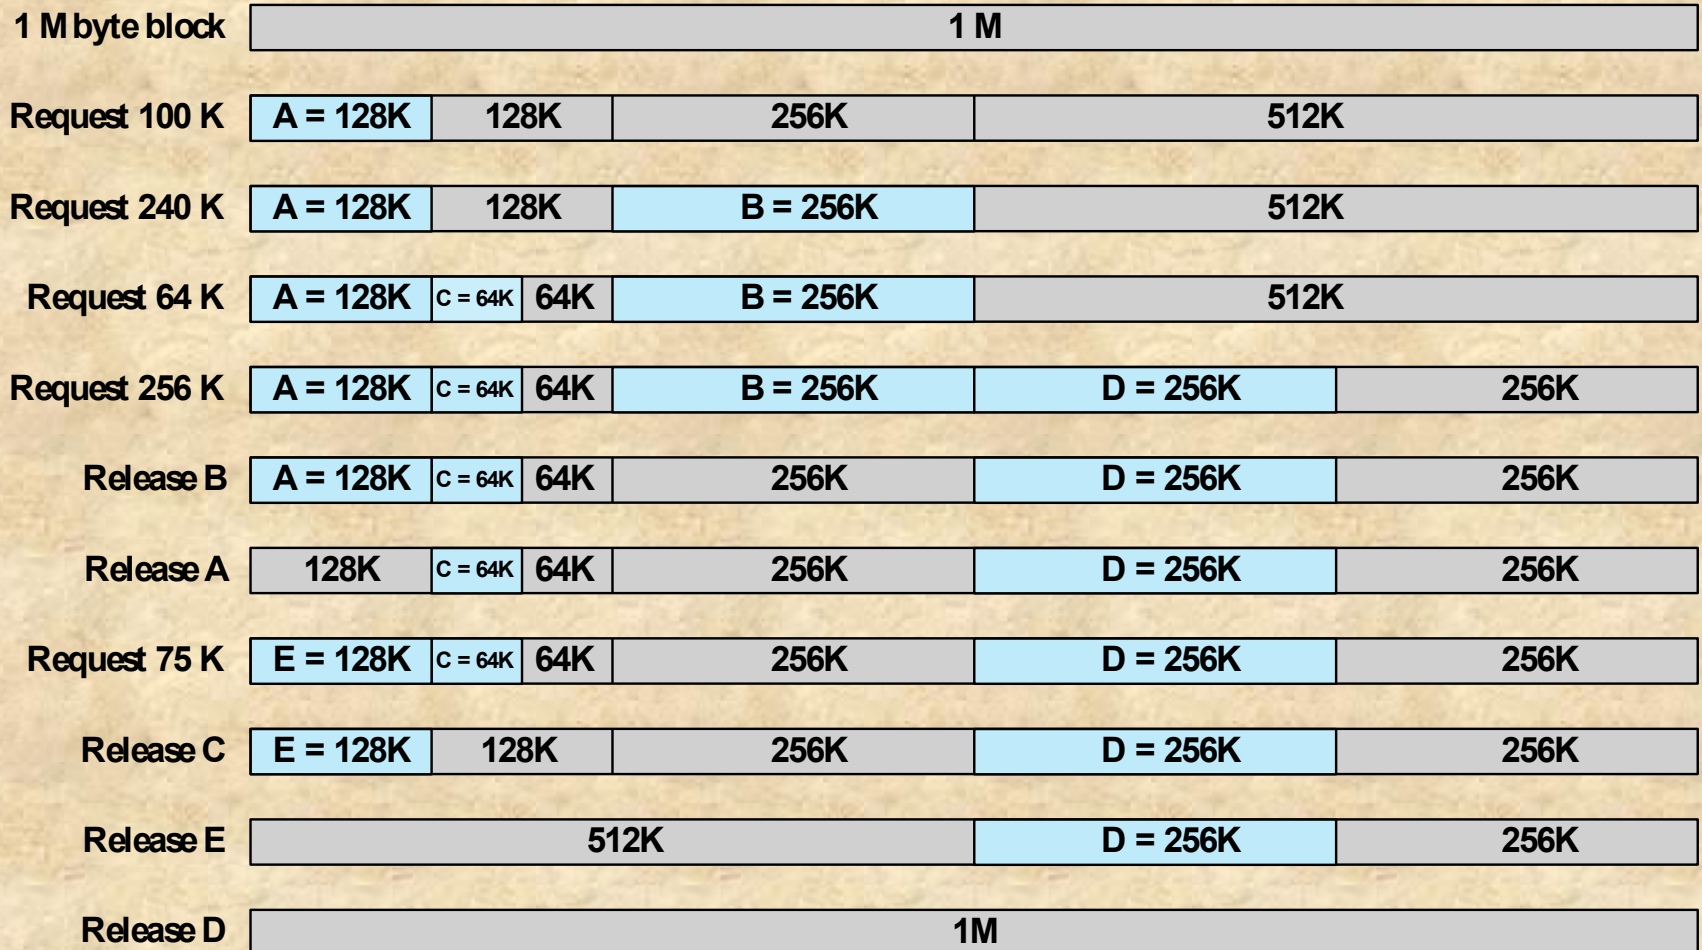
**Figure 7.5 Example Memory Configuration before and after Allocation of 16-M byte Block**



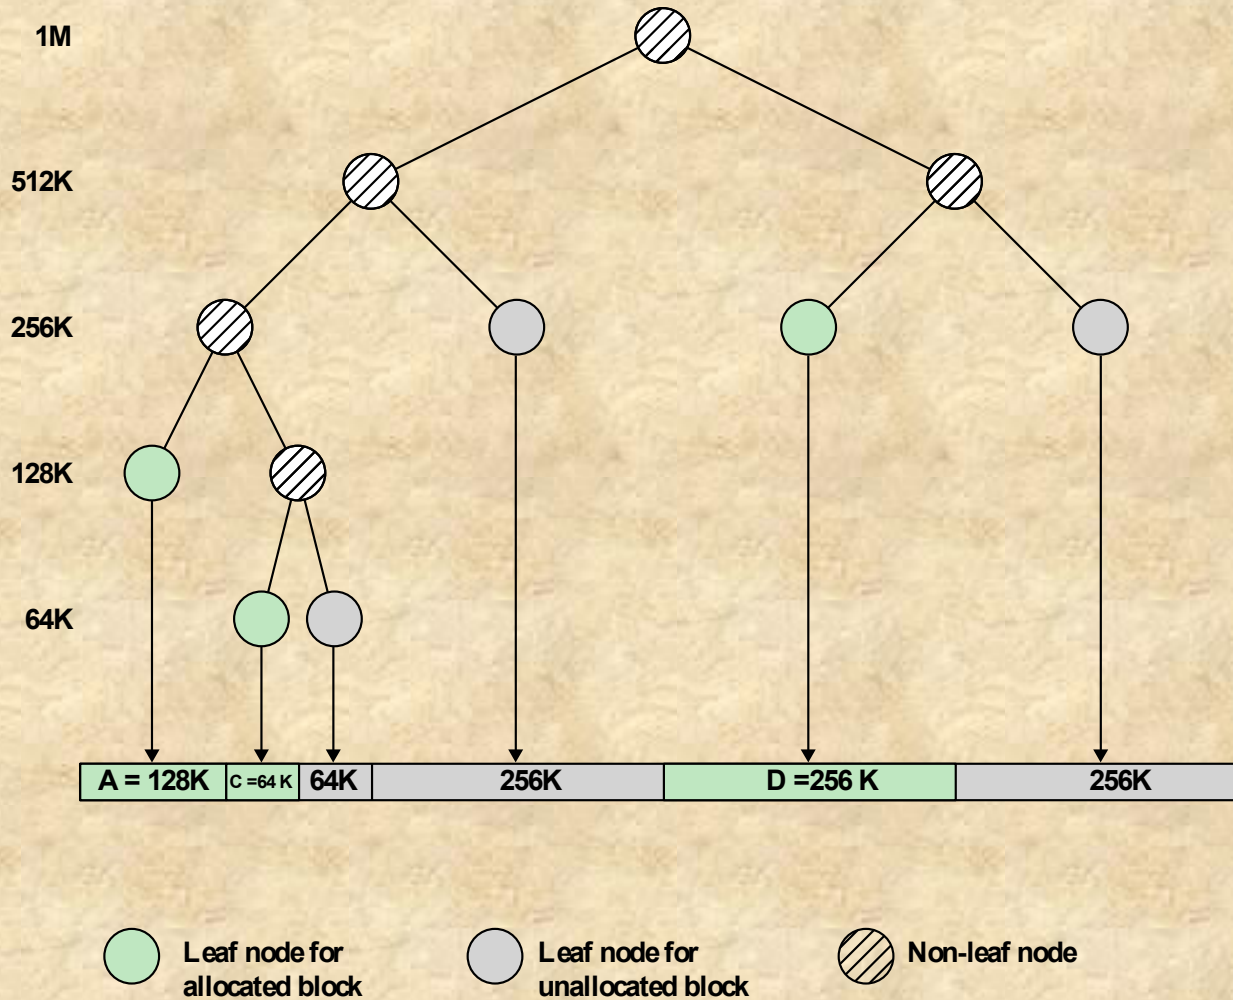
# Buddy System

- Comprised of fixed and dynamic partitioning schemes
- Space available for allocation is treated as a single block
- Memory blocks are available of size  $2^K$  words,  $L \leq K \leq U$ , where
  - $2^L$  = smallest size block that is allocated
  - $2^U$  = largest size block that is allocated; generally  $2^U$  is the size of the entire memory available for allocation





**Figure 7.6 Example of Buddy System**



**Figure 7.7 Tree Representation of Buddy System**



# Addresses

## Logical

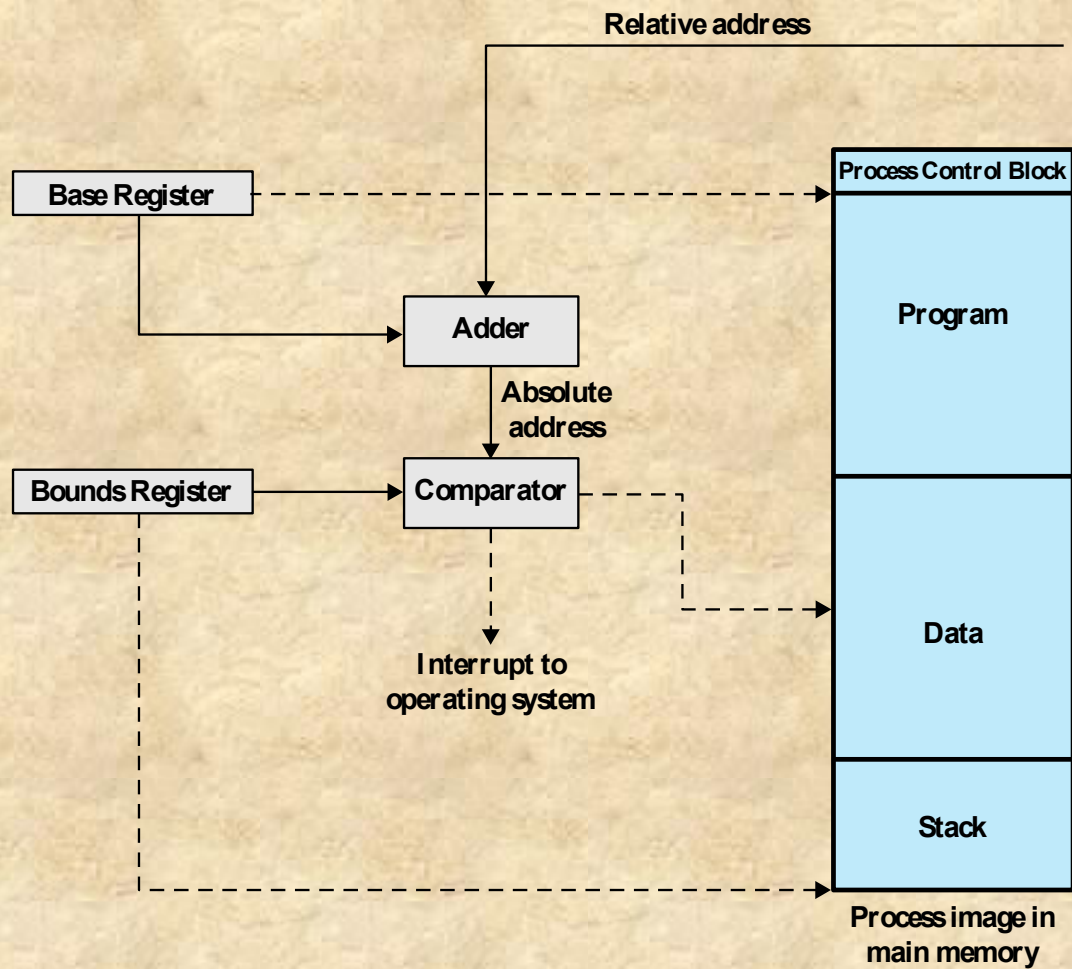
- reference to a memory location independent of the current assignment of data to memory

## Relative

- address is expressed as a location relative to some known point

## Physical or Absolute

- actual location in main memory



**Figure 7.8 Hardware Support for Relocation**

# Paging

- Partition memory into equal fixed-size chunks that are relatively small
- Process is also divided into small fixed-size chunks of the same size

## Pages

- chunks of a process

## Frames

- available chunks of memory



| Frame number | Main memory |
|--------------|-------------|
| 0            |             |
| 1            |             |
| 2            |             |
| 3            |             |
| 4            |             |
| 5            |             |
| 6            |             |
| 7            |             |
| 8            |             |
| 9            |             |
| 10           |             |
| 11           |             |
| 12           |             |
| 13           |             |
| 14           |             |

(a) Fifteen Available Frames

| Main memory |
|-------------|
| 0 A.0       |
| 1 A.1       |
| 2 A.2       |
| 3 A.3       |
| 4           |
| 5           |
| 6           |
| 7           |
| 8           |
| 9           |
| 10          |
| 11          |
| 12          |
| 13          |
| 14          |

(b) Load Process A

| Main memory |
|-------------|
| 0 A.0       |
| 1 A.1       |
| 2 A.2       |
| 3 A.3       |
| 4 B.0       |
| 5 B.1       |
| 6 B.2       |
| 7           |
| 8           |
| 9           |
| 10          |
| 11          |
| 12          |
| 13          |
| 14          |

(c) Load Process B

| Main memory |
|-------------|
| 0 A.0       |
| 1 A.1       |
| 2 A.2       |
| 3 A.3       |
| 4 B.0       |
| 5 B.1       |
| 6 B.2       |
| 7 C.0       |
| 8 C.1       |
| 9 C.2       |
| 10 C.3      |
| 11          |
| 12          |
| 13          |
| 14          |

(d) Load Process C

| Main memory |
|-------------|
| 0 A.0       |
| 1 A.1       |
| 2 A.2       |
| 3 A.3       |
| 4           |
| 5           |
| 6           |
| 7 C.0       |
| 8 C.1       |
| 9 C.2       |
| 10 C.3      |
| 11          |
| 12          |
| 13          |
| 14          |

(e) Swap out B


| Main memory |
|-------------|
| 0 A.0       |
| 1 A.1       |
| 2 A.2       |
| 3 A.3       |
| 4 D.0       |
| 5 D.1       |
| 6 D.2       |
| 7 C.0       |
| 8 C.1       |
| 9 C.2       |
| 10 C.3      |
| 11 D.3      |
| 12 D.4      |
| 13          |
| 14          |

(f) Load Process D

Figure 7.9 Assignment of Process Pages to Free Frames

# Page Table

- Maintained by operating system for each process
- Contains the frame location for each page in the process
- Processor must know how to access for the current process
- Used by processor to produce a physical address



|   |   |
|---|---|
| 0 | 0 |
| 1 | 1 |
| 2 | 2 |
| 3 | 3 |

**Process A  
page table**

|   |   |
|---|---|
| 0 | — |
| 1 | — |
| 2 | — |

**Process B  
page table**

|   |    |
|---|----|
| 0 | 7  |
| 1 | 8  |
| 2 | 9  |
| 3 | 10 |

**Process C  
page table**

|   |    |
|---|----|
| 0 | 4  |
| 1 | 5  |
| 2 | 6  |
| 3 | 11 |
| 4 | 12 |

**Process D  
page table**

|    |
|----|
| 13 |
| 14 |

**Free frame  
list**

**Figure 7.10 Data Structures for the Example of Figure 7.9 at Time Epoch (f)**



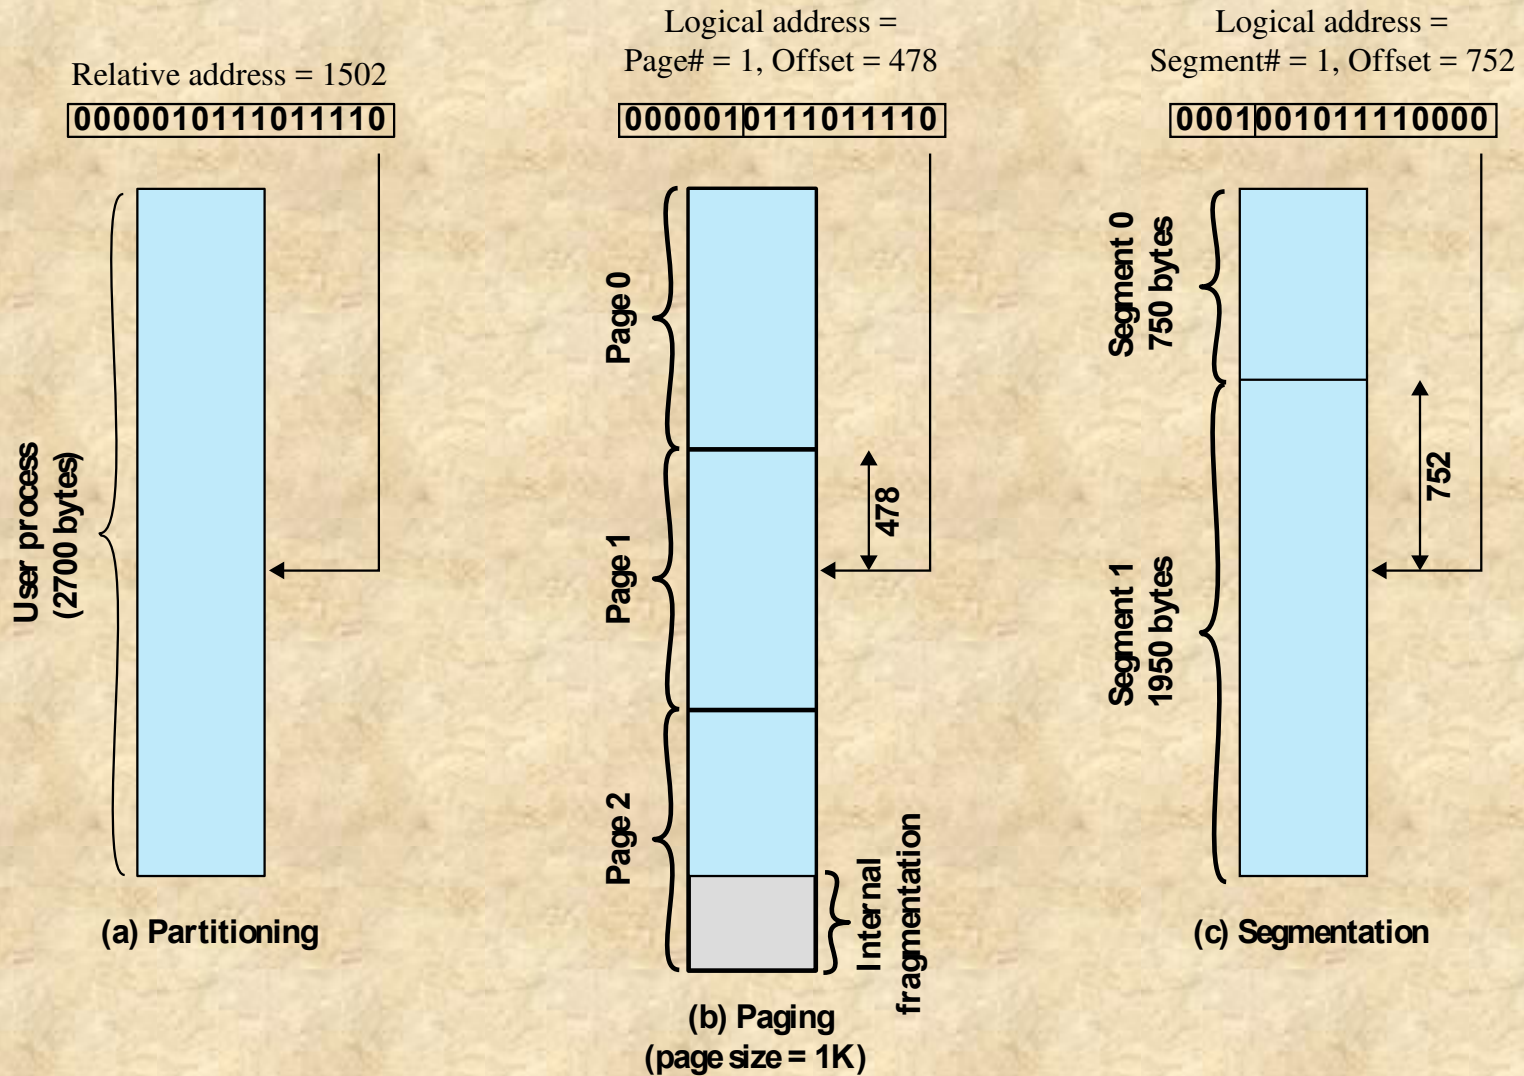
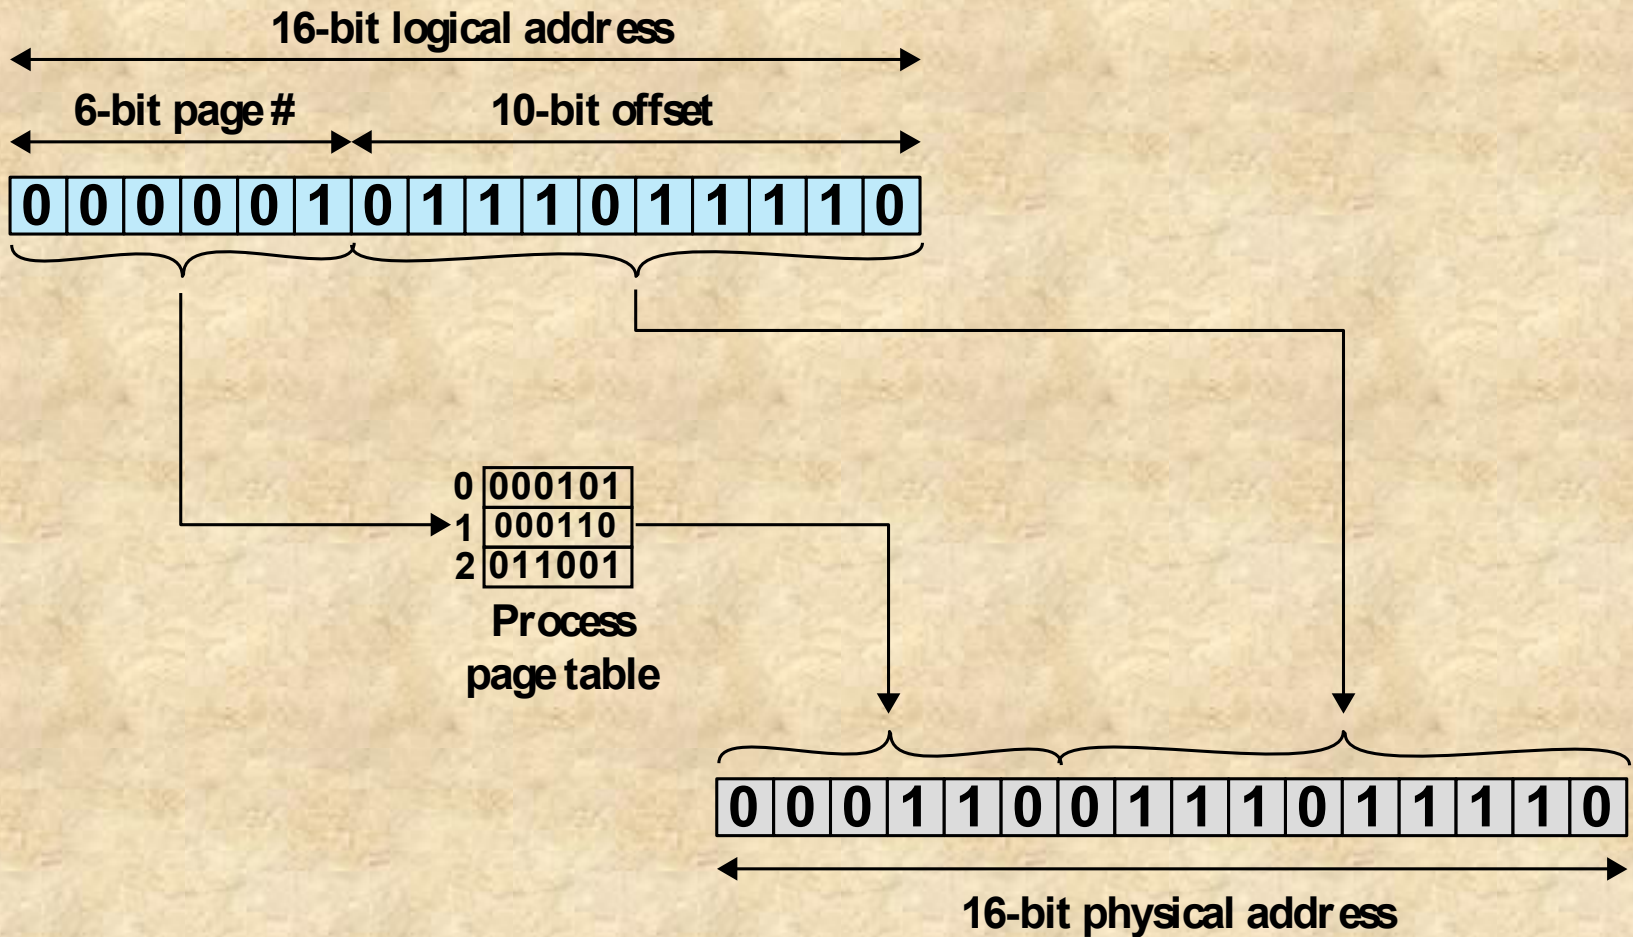


Figure 7.11 Logical Addresses



(a) Paging

Figure 7.12 Examples of Logical-to-Physical Address Translation

# Segmentation

- A program can be subdivided into segments
  - may vary in length
  - there is a maximum length
- Addressing consists of two parts:
  - segment number
  - an offset
- Similar to dynamic partitioning
- Eliminates internal fragmentation





# Segmentation

- Usually visible
- Provided as a convenience for organizing programs and data
- Typically the programmer will assign programs and data to different segments
- For purposes of modular programming the program or data may be further broken down into multiple segments
  - the principal inconvenience of this service is that the programmer must be aware of the maximum segment size limitation

# Address Translation

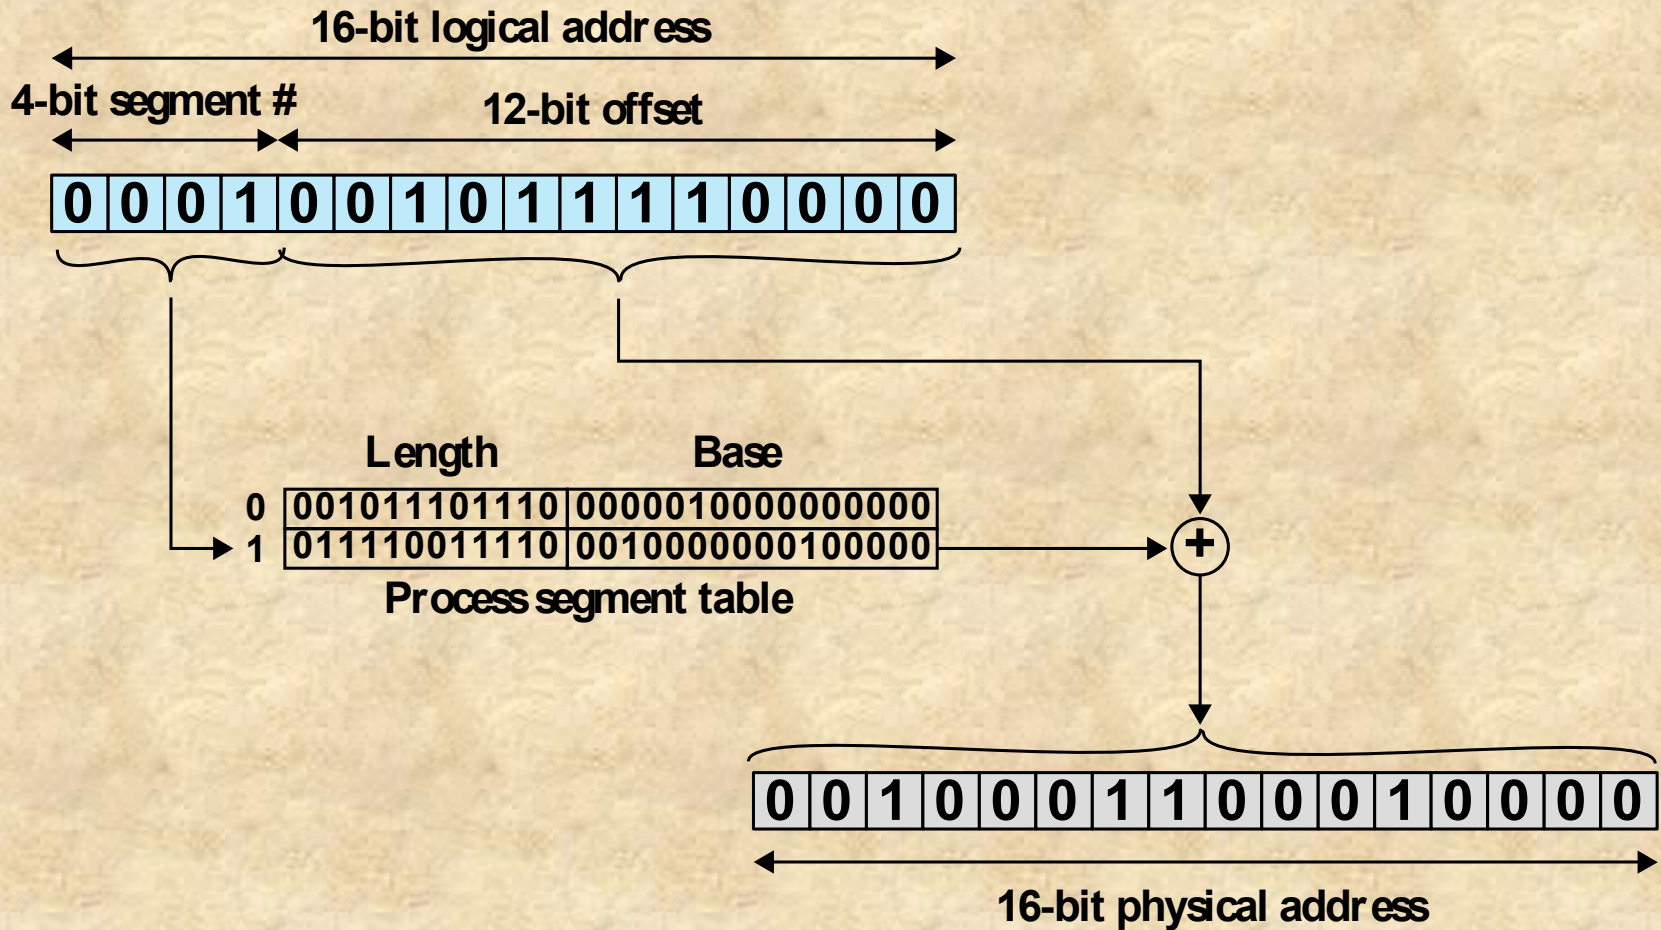
- Another consequence of unequal size segments is that there is no simple relationship between logical addresses and physical addresses
- The following steps are needed for address translation:

Extract the segment number as the leftmost  $n$  bits of the logical address

Use the segment number as an index into the process segment table to find the starting physical address of the segment

Compare the offset, expressed in the rightmost  $m$  bits, to the length of the segment. If the offset is greater than or equal to the length, the address is invalid

The desired physical address is the sum of the starting physical address of the segment plus the offset



(b) Segmentation

Figure 7.12 Examples of Logical-to-Physical Address Translation



# Summary

- Memory management requirements
  - relocation
  - protection
  - sharing
  - logical organization
  - physical organization
- Paging
- Memory partitioning
  - fixed partitioning
  - dynamic partitioning
  - buddy system
  - relocation
- Segmentation