


# COS 226


## Concurrency Chapter 1

### Acknowledgement

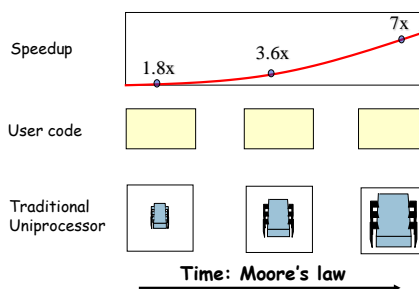


Some of the slides are taken from the companion slides for "The Art of Multiprocessor Programming" by Maurice Herlihy & Nir Shavit


### The Shared Memory Multiprocessor (SMP)




### Traditional Scaling Process



### Multicore Scaling Process



### Real-World Scaling Process




## Multiprocessor programming

- We look at concurrency from two directions:
  - Principles
    - Computability
  - Practice
    - Performance


## Model Summary

- Multiple threads
  - Sometimes called processes
- Single shared memory
- Objects live in memory
- Unpredictable asynchronous delays

### Mutual Exclusion or “Alice & Bob share a pond”



### Alice has a pet



### Bob has a pet



### The Problem



## Formalizing the problem

- First:
  - Mutual exclusion
  - This is a **safety** property – makes sure that nothing bad happens
- Both pets should never be in pond at the same time
  - Mutual exclusion
  - This is a **safety** property – makes sure that nothing bad happens

## And...

- If only one wants in, it gets in, but if both want in, only one gets in.
  - No deadlock
  - This is a **liveness** property – makes sure that something good happens eventually

## Simple Protocol

- A possible solution
  - Just look at the pond and see if the coast is clear
- Problem
  - Trees obscure the view

## Interpretation

- Threads can't "see" what other threads are doing
- Explicit communication required for coordination


## Cell Phone Protocol

- Another possible solution
  - Bob calls Alice (or vice-versa)
- Problem
  - Bob takes shower
  - Alice recharges battery
  - Bob out shopping for pet food ...

## Interpretation

- Message-passing doesn't work
- Recipient might not be
  - Listening
  - There at all
- Communication must be
  - Persistent (like writing)
  - Not transient (like speaking)

## Possible solution: Can Protocol




## Can Protocol


### ■ A possible solution:

- Bob puts one or more cans on Alice's windowsill attached to strings that lead to Bob's house
- When he wants to send a message he knocks over one of the cans
- When Alice sees the knocked over can, she resets them

## Bob conveys a bit



## Bob conveys a bit



## Can Protocol


### ■ Protocol

- Bob relies on Alice resetting the cans
- What if Alice goes away on holiday?
- Cans cannot be reused
- Bob runs out of cans

## Interpretation

- Cannot solve mutual exclusion with interrupts
  - Sender sets fixed bit in receiver's space
  - Receiver resets bit when ready
  - Requires infinite number of available bits


## Possible solution: Flag Protocol




## Alice's Protocol

- If Alice wants to release her pet she raises her flag
- If Bob's flag is down, she can release her pet
- When her pet returns, she lowers her flag again

## Alice's Protocol (sort of)



## Bob's Protocol (sort of)



## Bob's Protocol


- Raise flag
- When Alice's flag is down unleash pet
- Lower flag when pet returns

danger!

Bob's Protocol (2<sup>nd</sup> try)

- Raise flag
- While Alice's flag is up
  - Lower flag
  - Wait for Alice's flag to go down
  - Raise flag
- Unleash pet
- Lower flag when pet returns

## Bob's Protocol

- Raise flag
  - While Alice's flag is up
    - Lower flag
    - Wait for Alice's flag to go down
    - Raise flag
  - Unleash pet
  - Lower flag when pet returns
- Bob defers to Alice
- 

## The Flag Principle

- Raise the flag
- Look at other's flag
- Flag Principle:
  - If each raises and looks, then
  - Last to look must see both flags up

## Does it work?

- Mutual exclusion?
  - YES
  - Pets are not in the yard at the same time
- Deadlock-freedom?
  - YES
  - If both pets want to use the yard, Bob defers to Alice

## Starvation-freedom

- If a pet wants to enter the yard, will it eventually succeed?
  - NO.
  - Whenever Alice and Bob are in conflict, Bob defers to Alice, thus it is possible that Alice's pet uses the pond over and over and Bob's pet doesn't get a turn

## Waiting

- If Alice raises her flag and suddenly becomes ill, Bob's pet cannot use the pond until Alice returns
- Bob must *wait* for Alice to lower her flag

## Remarks

- Protocol is *unfair*
  - Bob's pet might never get in
- Protocol uses *waiting*
  - If Bob is eaten by his pet, Alice's pet might never get in


## The Fable Continues

- Alice and Bob fall in love & marry
- Then they fall out of love & divorce
  - She gets the pets – they now get along
  - He has to feed them – the pets however side with Alice and attacks Bob


## Producer-Consumer Problem

- A new coordination problem

## Bob Puts Food in the Pond



## Alice releases her pets to Feed



## Producer/Consumer

- Alice and Bob can't meet
  - Each has restraining order on other
  - So he puts food in the pond
  - And later, she releases the pets
- Avoid
  - Releasing pets when there's no food
  - Putting out food if uneaten food remains


## Producer/Consumer

- Need a mechanism so that
  - Bob lets Alice know when food has been put out
  - Alice lets Bob know when to put out more food


## Also known as bounded buffer problem

- Two processes – producer and consumer – share a common fixed-size buffer
- The producer generates data, puts it into the buffer and start again
- At the same time the consumer, consumes the data one piece at a time
- Problem:
  - Producer should not try to add data if the buffer is full
  - Consumer should not try to remove data from an empty buffer


## Surprise Solution



## Bob puts food in Pond




## Bob knocks over Can



## Alice Releases Pets



## Alice Resets Can when Pets are Fed






## Correctness

- Mutual Exclusion
  - Pets and Bob never together in pond
- No Starvation
  - if Bob always willing to feed, and pets always famished, then pets eat infinitely often.
- Producer/Consumer
  - The pets never enter pond unless there is food, and Bob never provides food if there is unconsumed food.

## Could Also Solve Using Flags




## Waiting

- Both solutions use waiting
- Waiting is **problematic**
  - If one participant is delayed
  - So is everyone else
  - But delays are common & unpredictable


## The Fable drags on ...

- Bob and Alice still have issues
- So they need to communicate
- So they agree to use billboards ...


## Billboards are Large




## Write One Letter at a Time ...




To post a message



Let's send another message



Uh-Oh



Readers/Writers

- Devise a protocol so that
  - Writer writes one letter at a time
  - Reader reads one letter at a time
  - Reader sees
    - Old message or new message
    - No mixed messages

Why do we care?

- Upgrading from a uniprocessor to a n-way multiprocessor does not mean in n-fold increase in performance
- We want as much of the code as possible to execute concurrently (in parallel)
- A larger sequential part implies reduced performance

Amdahl's law

- The extent to which we can speed up a complex job is limited by how much of the job must be executed sequentially.

## Amdahl's law

- Speedup = ratio between:
  - time it takes one processor to complete the task
    - Vs
  - time it takes  $n$  concurrent processors to complete the same task

## Amdahl's law

- $n$  – number of processors
- $p$  – fraction of task that can be executed in parallel
- Then:
  - The parallel part of the task will take  $p/n$  time
  - The sequential part of the task will take  $1 - p$  time
  - Parallelization is thus:  $1 - p + p/n$

## Amdahl's Law

$$\text{Speedup} = \frac{1}{1 - p + \frac{p}{n}}$$

## Example

- Ten processors
- 60% concurrent, 40% sequential
- How close to 10-fold speedup?

$$\text{Speedup} = 2.17 = \frac{1}{1 - 0.6 + \frac{0.6}{10}}$$

## Example

- Ten processors
- 80% concurrent, 20% sequential
- How close to 10-fold speedup?

$$\text{Speedup} = 3.57 = \frac{1}{1 - 0.8 + \frac{0.8}{10}}$$

## Example

- Ten processors
- 90% concurrent, 10% sequential
- How close to 10-fold speedup?

$$\text{Speedup} = 5.26 = \frac{1}{1 - 0.9 + \frac{0.9}{10}}$$

## The Moral

- Making good use of our multiple processors (cores) means
- Finding ways to effectively parallelize our code
  - Minimize sequential parts
  - Reduce idle time in which threads **wait** without

## Multicore Programming

- This is what this course is about...
  - The % that is not easy to make concurrent yet may have a large impact on overall speedup
- Next Week:
  - A more serious look at mutual exclusion