


COS 226

Chapter 2 Mutual Exclusion

Acknowledgement




- Some of the slides are taken from the companion slides for "The Art of Multiprocessor Programming" by Maurice Herlihy & Nir Shavit

Why is Concurrent Programming so Hard?

- Try preparing a seven-course banquet
 - By yourself
 - With one friend
 - With twenty-seven friends ...
- Before we can talk about programs
 - Need a language
 - Describing time and concurrency


Events

- An **event** a_0 of thread A is
 - Instantaneous
 - No simultaneous events (break ties)




Threads

- A **thread** A is (formally) a sequence a_0, a_1, \dots of events
 - "Trace" model
 - Notation: $a_0 \rightarrow a_1$ indicates order



Concurrency

- Thread A
- Thread B




Interleavings

- Events of two or more threads
 - Interleaved
 - Not necessarily independent




Intervals


- An *interval* $A_0 = (a_0, a_1)$ is
 - Time between events a_0 and a_1



Intervals may Overlap




Intervals may be Disjoint




Precedence

Interval A_0 *precedes* interval B_0




Precedence

- Notation: $A_0 \rightarrow B_0$
- Formally,
 - End event of A_0 before start event of B_0
 - Also called “happens before” or “precedes”



Precedence Ordering



- Remark: $A_0 \rightarrow B_0$ is just like saying
 - 1066 AD \rightarrow 1492 AD,
 - Middle Ages \rightarrow Renaissance,

Precedence Ordering



- Never true that $A \rightarrow A$
- If $A \rightarrow B$ then not true that $B \rightarrow A$
- If $A \rightarrow B$ & $B \rightarrow C$ then $A \rightarrow C$
- Funny thing: $A \rightarrow B$ & $B \rightarrow A$ might both be false!

Repeated Events

```
while (mumble) {
  a0; a1;
}
```

k -th occurrence
of event a_0

a_0^k

A_0^k

k -th occurrence of
interval $A_0 = (a_0, a_1)$

Implementing a Counter

```
public class Counter {
  private long value;

  public long getAndIncrement() {
    value++;
  }
}
```

Implementing a Counter

```
public class Counter {
  private long value;

  public long getAndIncrement() {
    temp = value;
    value = temp + 1;
    return value;
  }
}
```

Critical Section

- Block of code that can be executed by only one thread at a time
- Needs Mutual Exclusion
- Standard way to approach mutual exclusion is through locks

Locks (Mutual Exclusion)

```
public interface Lock {
    public void lock();
    public void unlock();
}
```

Locks (Mutual Exclusion)

```
public interface Lock {
    public void lock();
    public void unlock();
}
```

acquire lock

Locks (Mutual Exclusion)

```
public interface Lock {
    public void lock();
    public void unlock();
}
```

acquire lock

release lock

Using Locks

```
public class Counter {
    private long value;
    private Lock lock;
    public long getAndIncrement() {
        lock.lock();
        try {
            int temp = value;
            value = value + 1;
        } finally {
            lock.unlock();
        }
        return temp;
    }
}
```

Using Locks

```
public class Counter {
    private long value;
    private Lock lock;
    public long getAndIncrement() {
        lock.lock();
        try {
            int temp = value;
            value = value + 1;
        } finally {
            lock.unlock();
        }
        return temp;
    }
}
```

acquire Lock

Using Locks

```
public class Counter {
    private long value;
    private Lock lock;
    public long getAndIncrement() {
        lock.lock();
        try {
            int temp = value;
            value = value + 1;
        } finally {
            lock.unlock();
        }
        return temp;
    }
}
```

Release lock
(no matter what)

Using Locks

```
public class Counter {
    private long value;
    private Lock lock;
    public long getAndIncrement() {
        lock.lock();
        try {
            int temp = value;
            value = value + 1;
        } finally {
            lock.unlock();
        }
        return temp;
    }
}
```

Critical
section

Properties of a good Lock algorithm

- Mutual Exclusion
- Deadlock-free
- Starvation-free

Mutual Exclusion

- Threads do not access critical section at same time

Deadlock-free

- If some thread attempts to acquire the lock, some thread will succeed in acquiring the lock

Deadlock-free

- If **some** thread attempts to acquire the lock, **some** thread will succeed in acquiring the lock
 - System as a whole makes progress
 - Even if individuals starve
 - At least one thread is completing

Starvation-free

- Every thread that attempts to acquire the lock will eventually succeed

Starvation-free

- **Every** thread that attempts to acquire the lock will eventually succeed
 - If a thread calls lock() it will eventually acquire the lock
 - Individual threads make progress

Locks

- Let's start with lock solutions for 2 concurrent threads...

Two-Thread Conventions

```
class ... implements Lock {
    ...
    // thread-local index, 0 or 1
    public void lock() {
        int i = ThreadID.get();
        int j = 1 - i;
    }
    ...
}
```

Two-Thread Conventions

```
class ... implements Lock {
    ...
    // thread-local index, 0 or 1
    public void lock() {
        int i = ThreadID.get();
        int j = 1 - i;
    }
    ...
}
```

Henceforth: *i* is current thread, *j* is other thread

LockOne

- Basic idea:
 - Thread indicates interest in acquiring lock
 - Checks to see if other thread is currently in critical section
 - If true, waits until other thread finishes
 - If not, enters critical section

LockOne

```
class LockOne implements Lock {
    private boolean[] flag = new boolean[2];

    public void lock() {
        flag[i] = true;
        while (flag[j]) {}
    }

    public void unlock() {
        flag[i] = false;
    }
}
```

LockOne

```
class LockOne implements Lock {
    private boolean[] flag =
        new boolean[2];
    public void lock() {
        flag[i] = true;
        while (flag[j]) {}
    }
    public void unlock() {
        flag[i] = false;
    }
}
```

Set my flag

LockOne

```
class LockOne implements Lock {
    private boolean[] flag =
        new boolean[2];
    public void lock() {
        flag[i] = true;
        while (flag[j]) {}
    }
    public void unlock() {
        flag[i] = false;
    }
}
```

Set my flag

Wait for other flag to go false

LockOne

```
class LockOne implements Lock {
    private boolean[] flag =
        new boolean[2];
    public void lock() {
        flag[i] = true;
        while (flag[j]) {}
    }
    public void unlock() {
        flag[i] = false;
    }
}
```

When release lock set my flag again

Deadlock Freedom?

Concurrent execution:

```
flag[i] = true;   flag[j] = true;
while (flag[j]){} while (flag[i]){}

```

- If each thread sets its flag to true and waits for the other, they will wait forever
- No deadlock freedom

LockOne Summary

- LockOne offers mutual exclusion
- When accessed sequentially, LockOne works fine
- However with concurrent threads, LockOne is not Deadlock-free

LockTwo

Basic idea:

- When attempting to acquire lock, offer to be the victim that has to defer to other thread
- While current thread is the victim, wait until other thread becomes the victim
- When current thread no longer the victim, enter the critical section

LockTwo

```
public class LockTwo implements Lock {
    private int victim;
    public void lockO {
        victim = i;
        while (victim == i) {};
    }

    public void unlockO {}
}
```

LockTwo

```
public class LockTwo implements Lock {
    private int victim;
    public void lockO {
        victim = i;
        while (victim == i) {};
    }

    public void unlockO {}
}
```

Let other go first

LockTwo

```
public class LockTwo implements Lock {
    private int victim;
    public void lockO {
        victim = i;
        while (victim == i) {};
    }

    public void unlockO {}
}
```

Wait for permission

LockTwo

```
public class Lock2 implements Lock {
    private int victim;
    public void lockO {
        victim = i;
        while (victim == i) {};
    }

    public void unlockO {}
}
```

Nothing to do

LockTwo Claims

■ Satisfies mutual exclusion

- If thread *i* in CS
- Then `victim == i`
- Cannot be both 0 and 1

```
public void LockTwoO {
    victim = i;
    while (victim == i) {};
}
```

LockTwo Summary

- LockTwo offers Mutual Exclusion
- Works fine with concurrent threads
- However results in Deadlock with sequential threads
- LockOne and LockTwo thus complement each other

Peterson Lock

- Combine LockOne and LockTwo
 - Enable successful sequential access provided by LockOne
 - Enables successful concurrent access provided by LockTwo

Peterson Lock

- Basic idea:
 - Current thread indicates interest in acquiring the lock
 - Current thread offers to be the victim
 - If no interest from other thread and no longer the victim, then continue to critical section

Peterson's Algorithm

```
public void lock() {
    flag[i] = true;
    victim = i;
    while (flag[j] && victim == i) {};
}
public void unlock() {
    flag[i] = false;
}
```

Peterson's Algorithm

```
public void lock() {
    flag[i] = true;
    victim = i;
    while (flag[j] && victim == i) {};
}
public void unlock() {
    flag[i] = false;
}
```

Announce I'm interested

Peterson's Algorithm

```
public void lock() {
    flag[i] = true;
    victim = i;
    while (flag[j] && victim == i) {};
}
public void unlock() {
    flag[i] = false;
}
```

Announce I'm interested

Defer to other

Peterson's Algorithm

```
public void lock() {
    flag[i] = true;
    victim = i;
    while (flag[j] && victim == i) {};
}
public void unlock() {
    flag[i] = false;
}
```

Announce I'm interested

Defer to other

Wait while other interested & I'm the victim

Peterson's Algorithm

```

public void lock() {
    flag[i] = true;
    victim = i;
    while (flag[j] && victim == i) {};
}
public void unlock() {
    flag[i] = false;
}

```

Announce I'm interested
 Defer to other
 Wait while other interested & I'm the victim
 No longer interested

Mutual Exclusion

```

public void lock() {
    flag[i] = true;
    victim = i;
    while (flag[j] && victim == i) {};
}

```

- If thread 0 in critical section,
 - flag[0] = true,
 - victim = 1
- If thread 1 in critical section,
 - flag[1] = true,
 - victim = 0

Cannot both be true

Deadlock Free

```

public void lock() {
    ...
    while (flag[j] && victim == i) {};
}

```

- Thread blocked
 - only at **while** loop
 - Only if other's flag is true
 - only if it is the **victim**
- Solo: other's flag is false
- Both: one or the other must not be the victim

Starvation Free

- Thread *i* would be blocked only if *j* repeatedly re-enters so that

```

flag[j] == true and
victim == i

```

- When *j* re-enters
 - it sets **victim** to *j*.
 - So *i* gets in
- Thus: Starvation free

```

public void lock() {
    flag[i] = true;
    victim = i;
    while (flag[j] && victim == i) {};
}
public void unlock() {
    flag[i] = false;
}

```

Locks

- Moving on to solutions for *n* concurrent threads

Filter Lock

- Peterson lock adapted to work with *n* threads instead of just 2
- Thread has to traverse *n*-1 waiting rooms in order to acquire the lock

The Filter Algorithm for n Threads

There are $n-1$ “waiting rooms” called levels

- At each level
 - At least one enters level
 - At least one blocked if many try
- Only one thread makes it through



Filter

```
class Filter implements Lock {
    int[] level; // level[i] for thread i
    int[] victim; // victim[L] for level L

    public Filter(int n) {
        level = new int[n];
        victim = new int[n];
        for (int i = 1; i < n; i++) {
            level[i] = 0;
        }
    }
}
```

Filter

```
class Filter implements Lock {
    ...

    public void lock() {
        for (int L = 1; L < n; L++) {
            level[i] = L;
            victim[L] = i;
            while ((∃ k != i) level[k] >= L) &&
                victim[L] == i);
        }
    }

    public void unlock() {
        level[i] = 0;
    }
}
```

Filter

```
class Filter implements Lock {
    ...

    public void lock() {
        for (int L = 1; L < n; L++) {
            level[i] = L;
            victim[L] = i;
            while ((∃ k != i) level[k] >= L) &&
                victim[L] == i);
        }
    }

    public void release(int i) {
        level[i] = 0;
    }
}
```

One level at a time

Filter

```
class Filter implements Lock {
    ...

    public void lock() {
        for (int L = 1; L < n; L++) {
            level[i] = L;
            victim[L] = i;
            while ((∃ k != i) level[k] >= L) &&
                victim[L] == i); // busy wait
        }
    }

    public void release(int i) {
        level[i] = 0;
    }
}
```

Announce intention to enter level L

Filter

```
class Filter implements Lock {
    int level[n];
    int victim[n];
    public void lock() {
        for (int L = 1; L < n; L++) {
            level[i] = L;
            victim[L] = i;
            while ((∃ k != i) level[k] >= L) &&
                victim[L] == i);
        }
    }

    public void release(int i) {
        level[i] = 0;
    }
}
```

Give priority to anyone but me

Filter

Wait as long as someone else is at same or higher level, and I'm designated victim

```
class Filter implements Lock {
    int level[n];
    int victim[n];
    public void lock() {
        for (int L = 1; L < n; L++) {
            level[i] = L;
            victim[L] = i;
            while (( $\exists k \neq i$ ) level[k] >= L) &&
                victim[L] == i);
        }
    }
    public void release(int i) {
        level[i] = 0;
    }
}
```

Filter

```
class Filter implements Lock {
    int level[n];
    int victim[n];
    public void lock() {
        for (int L = 1; L < n; L++) {
            level[i] = L;
            victim[L] = i;
            while (( $\exists k \neq i$ ) level[k] >= L) &&
                victim[L] == i);
        }
    }
    public void release(int i) {
        level[i] = 0;
    }
}
```

Thread enters level L when it completes the loop

No Starvation

- Filter Lock satisfies properties:
 - Just like Peterson Algorithm at any level
 - So no one starves
- But what about fairness?
 - Threads can be overtaken by others

Waiting

- Starvation freedom guarantees that every thread that calls lock() eventually enters the critical section
- It however makes no guarantee about how long that can take

Waiting

- Ideally if A calls lock() before B, then A should enter critical section before B
- However this does not currently work since we cannot determine which thread called lock() first
- Locks should thus be further defined

Fairness

- Locks should be first-some-first served

Filter Lock again

- Filter Lock satisfies properties:
 - No one starves
 - But very weak fairness
 - Can be overtaken **arbitrary** # of times
 - That's pretty lame...

Bakery Algorithm

- Provides First-Come-First-Served
- How?
 - Take a "number"
 - Wait until lower numbers have been served
- Each thread takes a number when attempting to acquire the lock and waits until no thread with an earlier number is trying to acquire it

Bakery Algorithm

```
class Bakery implements Lock {
    boolean[] flag;
    Label[] label;
    public Bakery (int n) {
        flag = new boolean[n];
        label = new Label[n];
        for (int i = 0; i < n; i++) {
            flag[i] = false; label[i] = 0;
        }
    }
    ...
}
```

Bakery Algorithm

- flag[A] is a boolean flag indicating whether A wants to enter the critical section
- label[A] is an integer that contains thread A's "number" when entering the bakery

Bakery Algorithm

```
class Bakery implements Lock {
    ...
    public void lock() {
        flag[i] = true;
        label[i] = max(label[0], ..., label[n-1])+1;
        while (∃k flag[k]
            && (label[k] < label[i]));
    }
}
```

Bakery Algorithm

```
class Bakery implements Lock {
    ...
    public void lock() {
        flag[i] = true;
        label[i] = max(label[0], ..., label[n-1])+1;
        while (∃k flag[k]
            && (label[k] < label[i]));
    }
}
```

I'm interested

Bakery Algorithm

```

class Bakery implements Lock {
    ...
    public void lock() {
        flag[i] = true;
        label[i] = max(label[0], ..., label[n-1])+1;
        while (∃k flag[k]
                && (label[k] < label[i]));
    }
}

```

Take increasing label (read labels in some arbitrary order)

Label is created as one greater than the maximum of the other thread's labels

Bakery Algorithm

```

class Bakery implements Lock {
    ...
    public void lock() {
        flag[i] = true;
        label[i] = max(label[0], ..., label[n-1])+1;
        while (∃k flag[k]
                && (label[k] < label[i]));
    }
}

```

Someone is interested

Bakery Algorithm

```

class Bakery implements Lock {
    ...
    public void lock() {
        flag[i] = true;
        label[i] = max(label[0], ..., label[n-1])+1;
        while (∃k flag[k]
                && (label[k] < label[i]));
    }
}

```

Someone is interested

And someone has a smaller number than me

THEORETICALLY

Bakery Algorithm

- If two threads try to acquire the lock concurrently, they may read the same maximum number
- Threads thus have unique pairs consisting of number as well as thread ID

Bakery Algorithm

$$(label[i], i) << (label[j], j)$$

If and only if

$$label[i] < label[j] \text{ OR } label[i] = label[j] \text{ and } i < j$$

Bakery Algorithm

```

class Bakery implements Lock {
    boolean flag[n];
    int label[n];
    ...
    public void lock() {
        flag[i] = true;
        label[i] = max(label[0], ..., label[n-1])+1;
        while (∃k flag[k]
                && (label[k], k) << (label[i], i));
    }
}

```

Someone is interested

With lower (label, i) in lexicographic order

Bakery Algorithm

- In other words:
- Thread A must wait if:
 - Another thread is interested AND the other thread's number is lower than thread A
 - OR
 - Another thread is interested AND the two threads have the same number but the other's threads ID is smaller than A

Bakery Algorithm

```
class Bakery implements Lock {
    ...
    public void unlock() {
        flag[i] = false;
    }
}
```

Bakery Algorithm

```
class Bakery implements Lock {
    ...
    public void unlock() {
        flag[i] = false;
    }
}
```

No longer
interested

To analyse:

- Does the lock provide:
 - Mutual exclusion?
 - YES – two threads cannot be in the critical section at the same time since one of them will have an earlier label pair

To analyse:

- Starvation freedom?
 - YES – if a thread exists the critical section and immediately wants to reacquire the lock, he will first have to take a new, later number allowing the other waiting threads to gain access first

To analyse:

- Deadlock freedom?
 - YES – there is always one thread with the earliest label, ties are not possible because of labels consist of number and order in array

To analyse:

- The Bakery algorithm also provides First-come-first-served
 - If A calls lock() before B, then A's number is smaller than B's number
 - So B is locked out while flag[A] is true

Potential issue:

- With the current Bakery algorithm we are assuming that we have an infinite amount of numbers to use
- In practice this is not the case

Bounded timestamps

- Labels in the Bakery lock grow without bounds
- In a long-lived system we may have to worry about overflow
- If a thread's label silently rolled over from a large number to zero, the first-come-first-served property no longer holds

Bounded timestamps

- In the Bakery algorithm, the idea of labels can be replaced by timestamps
- Timestamps can ensure order among the contending threads
- We will thus need to ensure that if one thread takes a label after another, then the latter has the higher timestamp

Bounded timestamps

- Timestamps need the ability to:
 - Scan – read the other thread's timestamps
 - Label – assign itself a larger timestamp

Possible solution

- To construct a **Sequential** timestamping system
- Each thread perform scan-and-label completely one after the other
- Uses mutual exclusion

Bakery algorithm

- The Bakery algorithm is elegant and fair
- However it is not considered practical
 - Why?
 - Principal drawback is the need to read n distinct location where n can be very large