

Introduction to 64 Bit Intel Assembly Language Programming

Edited from the work of Ray Seyfarth

July 20, 2016

Goals for Cos 284

- Learn internal data formats
- Learn basic 64 bit Intel/AMD instructions
- Write pure assembly programs
- Write mixed C and assembly programs
- Use the gdb debugger for ASM
- Floating point instructions
- Arrays
- Functions
- Structs
- Using system calls and C libraries
- Data structures and high performance ASM

Problems with assembly language

- Assembly is the poster child for non-portability
 - ▶ Different CPU = different assembly
 - ▶ Different OS = different function ABI (application binary interface)
 - ▶ Intel/AMD CPUs operate in 16, 32 and 64 bit modes
- Difficult to program
 - ▶ More time = more money
 - ▶ Less reliable
 - ▶ Difficult to maintain
- Syntax does not resemble mathematics
- No syntactic protection
 - ▶ No structured ifs, loops
- No typed variables
 - ▶ Can use a pointer as a floating point number
 - ▶ Can load a 4 byte integer from a double variable
- Variable access is roughly like using pointers

What's good about assembly language?

- Assembly language is fast
 - ▶ Optimizing C/C++ compilers will be faster than a novice most of the time.
 - ▶ You need to dissect an algorithm and rearrange it to use a special feature that the compiler can't figure out
 - ▶ Generally you must use a special instructions
 - ▶ There are over 1000 instructions
 - ▶ Still it can be faster
- Assembly programs are small
 - ▶ But memory is cheap and plentiful
 - ▶ C/C++ compilers can optimize for size
 - ▶ Compilers can re-order code sections to reduce size
- Assembly can do things not possible in C/C++
 - ▶ I/O instructions
 - ▶ Manage memory mapping registers
 - ▶ Manipulate other internal control registers

What's good about assembly for ordinary mortals?

- Teaches you how the programs really works
- Teaches you how storage and arithmetic is done in registers
- Teaches you C/C++ function register and stack usages
- Teaches you how stack frames are built and destroyed.
- Optimization techniques are explained.
- Computer bugs are more immediately related to machine instructions and limitations
- You will learn how the compiler implements
 - ▶ if/else statements
 - ▶ loops
 - ▶ functions
 - ▶ structures
 - ▶ arrays
 - ▶ recursion
- Your coding will improve.

Generation of languages

- First generation - machine language
- Second generation - assembly language
 - ▶ Names for instructions
 - ▶ Names for variables
 - ▶ Names for locations of instructions
 - ▶ Perhaps with macros - code replacement
- Third generation - not machine instructions
 - ▶ Modeled after mathematics - Fortran
 - ▶ Modeled after English - Cobol
 - ▶ List processing - Lisp
- Fourth generation - domain specific
 - ▶ SQL
- Fifth generation - describe problem, computer generates algorithm
 - ▶ Prolog

Assembly example

```
; Program: exit
;
; Executes the exit system call
;
; No input
;
; Output: only the exit status ($? in the shell)
;
segment .text
global _start
_start:
    mov     eax,1          ; 1 is the exit syscall number
    mov     ebx,5          ; the status value to return
    int     0x80           ; execute a system call
```

Assembly syntax

- ; starts comments
- Labels are strings which are not instructions
 - ▶ Usually start in column 1
 - ▶ Can end with a colon to avoid confusion with instructions
- Instructions can be machine instructions or assembler instructions
 - ▶ `mov` and `int` are machine instructions or opcodes
 - ▶ `segment` and `global` are assembler instructions or pseudo-ops
- Instructions can have operands
 - ▶ `here: mov eax, 1`
 - ▶ `here` is a label for the instruction
 - ▶ `mov` is an opcode
 - ▶ `eax` and `1` are operands

Some assembler instructions

- `section` or `segment` define a part of the program
 - ▶ `.text` is where instructions go for Linux
- `global` defines a label to be used by the linker
- `global _start` makes `_start` a global label
- `_start` or `main` is where a program starts
 - ▶ `_start` is more basic
 - ▶ `main` is called (perhaps indirectly) by `_start`

Assembling the exit program

- `yasm -f elf64 -g dwarf2 -l exit.lst exit.asm`
- `-f elf64` says we want a 64 bit object file (elf=extensible linking format)
- `-g dwarf2` says we want dwarf2 debugging info (why dwarf?)
 - ▶ dwarf2 works pretty well with the gdb debugger
- `-l exit.lst` asks for a listing in `exit.lst`
- `yasm` will produce `exit.o`, an object file
 - ▶ machine instructions not ready to execute

exit.lst

```
1                                     %line 1+1 exit.asm
2
3
4
5
6
7
8
9
10                                  [segment .text]
11                                  [global _start]
12
13                                  _start:
14 00000000 B801000000               mov eax,1
15 00000005 BB05000000               mov ebx,5
16 0000000A CD80                     int 0x80
```

Linking

- Linking means combining object files to make an executable file
- For programs with `_start`
 - ▶ `ld -o exit exit.o`
 - ▶ Builds a file named `exit`
 - ▶ Default is `a.out`
- For programs with `main`
 - ▶ `gcc -o exit exit.o`
 - ▶ Gets default `_start` function from the C library
- `./exit` to run the program

Floating point numbers

Consider 1.75, in 32bit-IEEE 754 the number becomes:

0	01111111	110000000000000000000000
Positive	127	(1).75

Grouping into 4 bit nibbles:

0011	1111	1110	0000	0000	0000	0000	0000
3	f	e	0	0	0	0	0

But this is stored reversed and with each nibble pair swapped:

0	0	0	0	e	0	3	f
---	---	---	---	---	---	---	---

listings example part 1

Consider the following asm file “fp.asm”.

```
1      segment .data
2 zero  dd      0.0
3 one   dd      1.0
4 neg1  dd      -1.0
5 a     dd      1.75
6 b     dd      122.5
7 d     dd      1.1
8 e     dd      10000000000.0
```

The **dd** command specifies a double word data item. A word is 2 bytes.
So a double word is 32 bits.

- **dw** is a data word
- **db** is a byte
- **dq** is a data quad-word

listings example part 1

Now if we create the file listing using:

```
yasm -f elf64 -g dwarf2 -l fp.lst fp.asm
```

The result is:

1	%line 1+1 fp.asm
2	[section .data]
3 00000000 00000000	zero dd 0.0
4 00000004 0000803F	one dd 1.0
5 00000008 000080BF	neg1 dd -1.0
6 0000000C 0000E03F	a dd 1.75
7 00000010 0000F542	b dd 122.5
8 00000014 CDCC8C3F	d dd 1.1
9 00000018 F9021550	e dd 10000000000.0

Memory mapping

- Computer memory is an array of bytes from 0 to $n - 1$ where n is the memory size
- Programs perceive “logical” addresses which are mapped to physical addresses
- 2 people can run a program starting at logical address 0x4004c8 while using different physical memory
- CPU translates logical addresses to physical during instruction execution
- The CPU translation can be just as fast as if the software used physical addresses
- The x86-64 CPUs can map pages of sizes 4096 bytes and 2 megabytes
- Linux uses 2 MB pages for the kernel and 4 KB pages for programs
- Some recent CPUs support 1 GB pages

Translating an address

- Suppose an instruction references address 0x43215628
- With 4 KB pages, the rightmost 12 bits are an offset into a page
- With 0x43215628 the page offset is 0x628
- The page number is 0x43215
- Let's assume that the computer is set up to translate page 0x43215 to physical addresses 0x7893000 - 0x7893fff
- Then address 0x43215628 is mapped to 0x7893628

Benefits of memory mapping


- User processes are protected from each other
 - ▶ Your process can't read my process's data
 - ▶ Your process can't write my data
- The operating system is protected from malicious or errant code
- It is easy for the operating system to give processes contiguous chunks of “logical” memory

Why study memory mapping?

- If you write programs, the mapping is automatic
- We will not discuss instructions for changing mapping tables
- So what difference does it make?
- It helps explain page faults
 - ▶ Suppose you allocate an array of 256 bytes at logical address 0x45678200
 - ▶ Then all addresses from 0x45678000 to 0x45678fff are valid
 - ▶ You can go well past the end of the array before you can get a segmentation violation
- Knowledge is power!

Process memory model in Linux

- A Linux process has 4 logical segments
 - ▶ text: machine instructions
 - ▶ data: static data initialized when the program starts
 - ▶ heap: data allocated by `malloc` or `new`
 - ▶ stack: run-time stack
 - ★ return addresses
 - ★ some function parameters
 - ★ local variables for functions
 - ★ space for temporaries
- In reality it is more complex
- 141TB is 47 bits of all 1's
- CPU could use 48 bit logical addresses



Memory segments

- The text segment is named `.text` in `yasm`
 - ▶ `_start` and `main` are not actually at 0
 - ▶ The text segment does not need to grow, so the data segment can be placed immediately after it
- The data segment is in 2 parts
 - ▶ `.data` which contains initialized data
 - ▶ `.bss` which contains reserved data (initialized to 0)
 - ▶ “bss” stands for “Block Started by Symbol”
- The heap and the stack both need to grow
 - ▶ The heap grows up
 - ▶ The stack grows down
 - ▶ They meet in the middle and explode

Stack segment limits

- The stack segment is limited by the Linux kernel
- The typical size is 16 MB for 64 bit Linux
- This can be inspected using “`ulimit -a`” or “`ulimit -s`”
- 16 MB seems fairly small, but it is fine until you start using large arrays as local variables in functions
- The stack address range is 0x7fffffff000000 to 0xffffffffffff
- A fault to addresses in this range are recognized by the kernel to allow the stack to grow as needed

Memory example source code

```

        segment .data
a        dd      4
b        dd      4.4
c        times   10 dd 0
d        dw      1, 2
e        db      0xfb
f        db      "hello world", 0

        segment .bss
g        resd     1
h        resd     10
i        resb     100
```

Memory example source code (2)

```
segment .text
global  main      ; let the linker know about main
main:
    push    rbp      ; set up a stack frame for main
    mov     rbp, rsp ; set rbp to point to the stack frame
    sub     rsp, 16   ; leave some room for local variables
                    ; leave rsp on a 16 byte boundary
    xor     eax, eax  ; set rax to 0 for return value
    leave   ; undo the stack frame manipulations
    ret
```


Memory example listing file

```
1                                %line 1+1 memory.asm
2                                [section .data]
3 00000000 04000000             a dd 4
4 00000004 CDCC8C40             b dd 4.4
5 00000008 00000000<rept>      c times 10 dd 0
6 00000030 01000200             d dw 1, 2
7 00000034 FB                   e db 0xfb
8 00000035 68656C6C6F20776F72-  f db "hello world", 0
9 00000035 6C6400
```

- Addresses are relative to start of .data in this file
- Notice that the 4 byte of 4 is at address 0 (backwards)
- $b = 0x408ccccd = 0\ 10000001\ 00011001100110011001101$
- Sign bit is 0, exponent field is $0x81 = 129$, exponent = 2
- Fraction is $1.00011001100110011001101$

Memory example listing file (2)

```
11                                [section .bss]
12 00000000 <gap>                g resd 1
13 00000004 <gap>                h resd 10
14 0000002C <gap>                i resb 100
```

- Notice that the addresses start again at 0
- The commands reserve space
- `resd 1` reserves 1 double word or 4 bytes
- `resd 10` reserves 10 double words or 40 bytes
- `resb 100` reserves 100 bytes

Memory example listing file (3)

16		[section .text]
17		[global main]
18		main:
19	00000000 55	push rbp
20	00000001 4889E5	mov rbp, rsp
21	00000004 4883EC10	sub rsp, 16
22	00000008 31C0	xor eax, eax
23	0000000A C9	leave
24	0000000B C3	ret

Examining memory

Useful tools to examine memory are:

- gdb
- ebe